

# Block Diffusion: Interpolating Between Autoregressive and Diffusion Language Models

Ntountounakis Georgios, Markoulidakis Georgios, Vitalis Petros,  
Makras Ilias, Kritharidis Konstantinos, Kordas Nikolaos

Pattern Recognition, ECE  
National Technical University of Athens

January 2026

# Table of Contents

1 Paper Overview

2 Implementation Plan

3 Computational Resources & Data

4 Team Organization

# Table of Contents

1 Paper Overview

2 Implementation Plan

3 Computational Resources & Data

4 Team Organization

# Introduction to the Problem-Motivation

**Two main approaches for Language Models:**

**Autoregressive (AR):**

- Token-by-token generation
- High quality
- KV caching
- Variable length

**Diffusion:**

- Parallel generation
- Better controllability
- **Fixed length (limitation)**
- **Lower quality (Perplexity Gap)**

**Question**

Can we combine the advantages of both approaches?

# Core Idea: Block Diffusion

## Key Innovation

Split sequence into **blocks** and apply:

- Autoregression **between** blocks
- Diffusion **within** each block



Diffusion within block(parallel)

Autoregressive over blocks

**Parameterization:** Block size  $L'$  controls the trade-off

- $L' = 1 \rightarrow$  Pure AR
- $L' = L \rightarrow$  Pure Diffusion

# Block Diffusion Training Algorithm

**Challenge:** To compute the loss  $\mathcal{L}_{BD}$  for block  $b$ :

- **Noisy Pass:** The model must process  $\mathbf{x}_t^b$  to learn denoising.
- **Clean Pass:** The model needs the clean  $\mathbf{x}^b$  to provide context for subsequent blocks.

**Naive Solution:** Every block must pass through  $\mathbf{x}_\theta$  at least *twice*.

---

## Algorithm 1 Block Diffusion Training

**Input:**  $\mathbf{x}, B, q_t, \mathbf{x}_\theta, \mathcal{L}_{BD}$

- 1: **repeat**
- 2:    Sample  $t_1, \dots, t_B \sim U[0, 1]$
- 3:     $\mathbf{x}_{t_b}^b \sim q_{t_b}(\cdot | \mathbf{x}^b)$
- 4:     $\emptyset, K_{1:B}, V_{1:B} \leftarrow \mathbf{x}_\theta(\mathbf{x}) \quad \triangleright \text{Clean Pass}$
- 5:     $\mathbf{x}_{\text{logit}}^b \leftarrow \mathbf{x}_\theta^b(\mathbf{x}_{t_b}^b, K_{1:b-1}, V_{1:b-1})$   
     $\triangleright \text{Noisy Pass}$
- 6:    Grad step  $\nabla_\theta \mathcal{L}_{BD}$
- 7: **until** converged

## Why this design?

- **Indep.  $t_b$ :** Reduces variance
- **KV caching:**  $5\times$  speedup
- **Vectorized:** FlexAttention parallelism

# Block Diffusion Sampling Algorithm

**Challenge:** Standard diffusion models cannot generate variable-length sequences.

---

## Algorithm 2 Block Diffusion Sampling

---

**Input:** # blocks  $B$ , model  $x_\theta$ , diffusion sampling algorithm SAMPLE

- 1:  $x, K, V \leftarrow \emptyset$      $\triangleright$  Initialize output & cache
  - 2: **for**  $b = 1$  to  $B$  **do**
  - 3:     $x^b \leftarrow \text{SAMPLE}(x_\theta^b, K_{1:b-1}, V_{1:b-1})$   
       $\triangleright$  Parallel block production
  - 4:     $(K, V) \leftarrow (K \oplus K_b, V \oplus V_b)$      $\triangleright$   
      Extend cache for new block
  - 5: **end for**
  - 6: **return**  $x$
- 

## Why this design?

- **Autoregressive over blocks:** Generate arbitrary length sequences
- **Diffusion within blocks:** Parallel token generation  
 $\rightarrow$  faster text generation
- **KV caching:** Reuse computation from previous blocks  $\rightarrow$  inference efficiency
- **Result:**  $40\times$  fewer steps than SSD-LM with better quality



# The Problem: Perplexity Gap

**Question:** Why do diffusion models have worse perplexity?

## Observation

For  $L' = 1$  (block size = 1), the diffusion objective is equivalent to the AR objective in expectation.

**But in practice:**

Model	PPL ( $\downarrow$ )
AR	<b>22.88</b>
+ random batch size	24.37
BD3-LM $L' = 1$	$\leq 25.56$
+ tuned schedule	<b>22.88</b>

Masking  $\Rightarrow$  diffusion training estimates loss gradients with 2x fewer tokens.

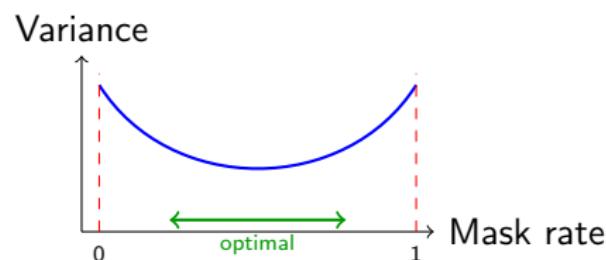
**Conclusion:** The problem is not the objective, but the **variance!**

# The Solution: Clipped Noise Schedules

**Goal:** Design schedules that minimize gradient variance (Section 5)

**Why extreme mask rates are bad:**

- Too few masked (near 0): Easy task, no learning signal
- All masked (near 1): Predict marginals, doesn't help generation
- Both cause **high variance gradients**



# The Solution: Clipped Noise Schedules

## Clipped Schedules

$$1 - \alpha_t \sim U[\beta, \omega]$$

- $\beta$ : Min mask rate,  $\omega$ : Max mask rate
- Avoids extremes near 0 or 1

## Examples:

- $U[0, 1]$ : Standard linear
- $U[0.3, 0.8]$ : Clipped
- $U[0.45, 0.95]$ : Heavy masking (best for  $L' = 4$ )

## Key Finding

There exists a unique distribution for each block size, that minimizes both the variance of the NELBO and the PPL.

# Main Experimental Results

**Table:** Test perplexities (PPL;  $\downarrow$ ) on OWT for models trained for 524B tokens. Best diffusion value is bolded.

PPL ( $\downarrow$ )	
AR	17.54
SEDD	$\leq$ 24.10
MDLM	$\leq$ 22.98
BD3-LMs $L' = 16$	$\leq$ 22.27
$L' = 8$	$\leq$ 21.68
$L' = 4$	$\leq$ <b>20.73</b>

**Table:** Generation length statistics from sampling 500 documents from models trained on OWT.

	Median # tokens	Max # tokens
OWT train set	717	131K
AR	4008	131K
SEDD	1021	1024
BD3-LM $L' = 16$	798	9982

# Table of Contents

1 Paper Overview

2 Implementation Plan

3 Computational Resources & Data

4 Team Organization

# Implementation Phases

Phase	Description	Weeks
1	Paper & codebase study	1-2
2	Environment & data setup	2-3
3	Initial reproduction	3-5
4	Final reproduction	5-6
5	Extension experiments	6-7
6	Analysis & Writing	7-8

## Current Status

Phase 3 (Initial reproduction): running a small-scale setup (tiny model + reduced dataset) to validate the full pipeline.

# Reproduction

**Goal:** Reproduce the paper's training/inference pipeline in a *controlled* setting.

## Training datasets and tokenizers

Training dataset	Tokenizer
LM1B	bert-base-uncased
OWT	GPT2

# Reproduction (experimental setup)

## Experimental Setup (Small-scale)

Configuration	Choice
Training Samples	500
Training Steps	800
Evaluation Samples	100
Model Size	tiny
Context Length	128

## What we will report:

- Training loss curves and evaluation perplexity (PPL)
- Sampling sanity checks: coherence vs. speed (block size trade-off)

**Expected outcome:** correct trends and stable training (not necessarily paper-level absolute PPL due to limited resources)

# Potential Extensions

## Extension 1: Diffusion Optimization

- Optimize elements of the diffusion process to improve quality and/or performance.

## Extension 2: Adaptive Block Size

- Adaptively change the block size according to the need for fine-grained detail of each block.

## Extension 3: Training / Fine-tuning on new, resource-friendly datasets

- **AG News**: news-domain text for domain adaptation
- **LAMBADA**: long-range dependency evaluation
- **Penn Treebank (ptb)**: classic language modeling benchmark
- **BabyLM Corpus**: child-language data for acquisition modeling

# Experimental Results (Reproduction)

Table 3 Initial Reproduction: Test Perplexities

PPL ( $\downarrow$ )	
<b>Autoregressive</b>	
Transformer	1221
<b>Diffusion</b>	
SEDD	$\leq 1403$
MDLM	$\leq 1370$
<b>Block diffusion (Ours)</b>	
BD3-LMs $L' = 16$	$\leq 1345$
$L' = 8$	$\leq 1210$
$L' = 4$	<b><math>\leq 1176</math></b>

# Table of Contents

1 Paper Overview

2 Implementation Plan

3 Computational Resources & Data

4 Team Organization

# Computational Resources & Data

## Computational Resources:

- Google Colab (Primary environment)
- Kaggle Notebooks
- our GPUs

## Datasets:

- LM1B
- OWT

# Table of Contents

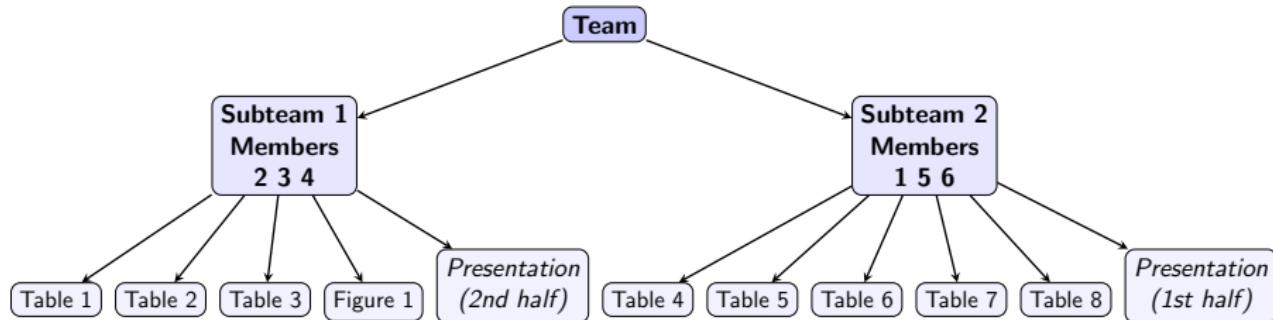
1 Paper Overview

2 Implementation Plan

3 Computational Resources & Data

4 Team Organization

# Team Organization and Structure



## Coordination:

- Weekly team meetings
- Subteams internal coordination
- GitHub repo
- Discord server

# Member Roles

Phase	Description	Members
1	Paper & codebase study	Everyone
2	Environment & data setup	Georgios M., Petros
3	Initial reproduction	Georgios M., Petros, Ilias
4	Final reproduction	Georgios Nt., Konstantinos, Nikolaos
5	Extension experiments	Everyone
6	Analysis & Writing	Everyone

Thank You!

Questions?

**Paper:** Block Diffusion (ICLR 2025)

**Code:** <https://github.com/kuleshov-group/bd3lms>

**Our Repo:** <https://github.com/ntua-el21050/bd3lms>

**Authors:** Arriola, Gokaslan, Chiu, Yang, Qi, Han, Sahoo, Kuleshov